

Collaborative Planning Product Context

Fictional product-design challenge dataset. Not real Apple internal information.

This briefing is the onboarding pack for the collaborative trip planning exploration inside the maps product. It was prepared by Nina (Product Manager), with contributions from Maya (UX Researcher), Leo (Interaction Designer), Alex (Content Designer) and Sam (Engineering), for a product designer joining mid-stream: it collects the brief, the research so far, the behavioural data we hold, the assumptions we are carrying, and the questions we have not answered. Read it alongside the primary research artefacts — the journey map, the workshop output, the survey, and the first depth interview — rather than instead of them. This pack summarises context; the artefacts carry the evidence.

1. Project brief

The maps product has historically treated saving places as a private, single-player activity. A person finds a restaurant, taps save, and the place lands in a personal list that nobody else can see. That model serves individual use well, but not the moment when a person is planning not for themselves but for a group — a long weekend with university friends, a family holiday with three generations, two couples coordinating a week away.

This exploration asks a deliberately open question: **what would it take for a small group of people to plan a trip together, from first idea through to travelling, with the maps product as a genuinely useful part of that process?**

The exploration is framed around a six-stage journey that recurs throughout this pack and the research artefacts:

Discover → Share → Save → Compare → Decide → Navigate

- **Discover:** an individual encounters a place — social video, a recommendation, a search, walking past it.
- **Share:** they surface it to the group, almost always as a link or screenshot in a messaging app.
- **Save:** someone (often one specific someone) records it somewhere more durable — a note, a list, a spreadsheet.
- **Compare:** the group weighs options against each other — price, location, opening hours, vibe.
- **Decide:** the group commits to specific places, dates and bookings.
- **Navigate:** on the trip itself, the plan is used in the world — directions, timings, changes on the fly.

The exploration is scoped to leisure trips planned by informal groups of roughly two to ten people, excluding business travel, large-event coordination and single-traveller planning, though findings may generalise.

The team is five people plus you. Nina owns the product framing and this brief. Maya has run all fieldwork and owns the research plan. Leo produced the journey map, co-facilitated the ideation workshop, and built the first interactive concept, working title **Trip Spaces**, now through one moderated evaluation session. Alex owns naming and the language of any shared surface. Sam is the engineering counterpart, assessing feasibility of shared-list infrastructure.

The exploration runs twelve weeks; this pack is issued at the start of week five. Expected outputs: a synthesised point of view, two to three concept directions with rationale, one developed to a testable prototype, and a recommendation on how to proceed. Your immediate task is to absorb the evidence, form your own view of where the real problem sits, and pressure-test Trip Spaces against that view rather than inheriting it uncritically.

2. Why collaborative planning matters

All figures in this section are directional category observations, not measurements of our own product.

Most trips are group trips. Category research consistently indicates that most leisure trips involve two or more travellers, and a substantial share involve three or more. Yet almost every planning tool in wide use, our own saved-places feature included, is built for one person's head. That mismatch is the founding observation of this exploration.

Planning is where the emotional investment lives. Diary studies across the travel category repeatedly find that anticipation and planning account for a large share of a trip's reported enjoyment — people describe the weeks of collecting ideas with as much warmth as the trip itself. A product present only at the Navigate stage arrives after most of those moments have happened elsewhere.

The maps product already holds half the raw material. People already save places, and place data — hours, price bands, photos, locations — is the substrate every group plan is made of. Today that substrate is scattered across screenshots and pasted links that go stale. There is a plausible thesis that bringing the group's planning to where the place data lives, rather than exporting fragments to wherever the group talks, produces a step change in usefulness.

Retention logic. A personal saved list is used by one person, occasionally. A shared trip plan is a surface several people return to over weeks, generating natural re-engagement moments and culminating in high-value actions — navigation sessions, bookings via place pages. If collaborative planning works, it converts a single-player utility into a repeated multi-player one. Whether the design case supports the business case is what this exploration exists to find out.

3. Current-state behaviour (fictional)

The following composite describes how the groups we have studied plan trips today. Drawn from Maya's fieldwork, it is deliberately descriptive: it records what people do, not what any of it means.

A trip typically begins in a messaging thread. Someone floats an idea — "we should do a weekend away before Christmas" — and the thread briefly ignites. Within a day or two, links start arriving: a social video of a food market, a flat listing, a screenshot of a viewpoint with no caption. These interleave with the group's normal conversation, so a link posted on Tuesday is four hundred messages up by Friday.

At some point one person — almost always the same person, trip after trip — begins consolidating: scrolling back through the thread, retyping places into a notes app, a spreadsheet, or occasionally a shared document. This artefact becomes the plan's de facto home, though it is disconnected from the place data itself: hours, prices and locations must be re-checked by hand, and the note goes stale the moment anything changes.

Comparison happens through improvisation: screenshots of maps with pins, two accommodation listings pasted side by side with "thoughts?", occasionally a spreadsheet with price and distance columns. The maps

app appears mostly as a lookup tool, and its saved lists are used privately, if at all — several participants described maintaining a personal saved list for a group trip that no one else in the group ever saw.

Decision-making is the least externally visible stage, happening in the chat, on calls, or in person, with little tooling involved. Booking happens on airline, rail, accommodation and restaurant sites — never inside the maps product — and confirmation screenshots are posted back into the thread, where they join everything else in the scroll.

On the trip, the plan fragments again: the organiser has the master note, others have partial memories of the thread, and the maps app finally gets heavy use for navigation — usually from fresh searches rather than from anything saved during planning.

The tools involved, in rough order of prominence:

Tool (generic)	Role in planning	Journey stages where it appears
Group chat in a messaging app	Default venue: ideas, links, reactions, decisions	All six stages
Notes app	Consolidated list, usually maintained by one person	Save, Compare, Decide
Social video apps	Primary discovery source for younger groups	Discover
Browser tabs	Held open for days as an informal shortlist	Discover, Compare
Spreadsheet	Structured comparison, in a minority of groups	Compare, Decide
Maps app (saved lists)	Personal saving; lookup during comparison; navigation on the trip	Save, Compare, Navigate
Shared document	Occasional plan-of-record for organised groups	Save, Compare, Decide
Screenshots	Universal transport format between all of the above	Share, Compare

Two recurring artefacts appear across nearly every group studied: the screenshot as a unit of currency (places, listings and confirmations travel as images, shedding their links and metadata in transit), and the versioned note ("Trip plan FINAL", "Trip plan FINAL v2") maintained by the organiser.

4. Research completed so far

The exploration has completed four primary research activities, each with a dedicated artefact you should read in full; these summaries cover method and coverage only.

Current-state journey map (Leo and Maya, weeks 1–3). A synthesis wall of roughly fifty observation stickies organised against the six-stage journey, built from workshop warm-ups, survey free text and ad-hoc conversations. It records where friction, workarounds and emotional peaks were observed — the team's shared map of where things happen, deliberately stopping short of explaining why.

Collaborative planning workshop (facilitated by Leo, week 3). A ninety-minute ideation session with the core team plus four adjacent colleagues, producing around forty idea and question notes against the journey stages. The output reflects the room's instincts as much as the evidence; read it as a catalogue of candidate directions, not a prioritised roadmap.

Group planning survey (Maya, weeks 2–4, n=40). A screener-plus-attitudes survey of adults who had planned or joined a group leisure trip within the past eighteen months, covering group composition, tools, organiser dynamics, sharing frequency, decision methods, satisfaction and free-text frustrations. Forty

respondents is a directional sample — large enough to see clusters, too small for statistical confidence — skewed towards UK respondents aged 22–45. Column definitions are in the CSV header notes.

Depth interview 01 with prototype evaluation (Maya, week 4). A single moderated session with one participant — a 27-year-old default organiser of a six-person university friend group — covering three recent or current trips, followed by a four-task walkthrough of the Trip Spaces concept. One voice, and also the richest behavioural material we hold; the transcript rewards close reading, including where what the participant says and what he does diverge. Further interviews are planned but not yet scheduled.

Desk research (ongoing). Generic category scanning of collaborative planning patterns; summarised in section 18.

What we have **not** done matters as much: no diary study of a live group planning end-to-end; no session observing a group deciding together (all our data is individuals reporting on their groups); no behavioural experiment. The telemetry in section 12 is observational and cannot speak to causation. Where this pack states something confidently, check which of the above it rests on.

5. User groups

Group travel is not one behaviour. The research surfaces recurring archetypes — of groups, and of roles within groups. These are working sketches, not validated segments, useful for keeping concept work honest about who it serves and who it quietly ignores.

Group archetypes

Archetype	Typical size	Cadence	Planning texture
University friend group	4–8	1–2 trips/year	Scattered cities, one default organiser, chat-centric, long collecting phase
Couple + couple	4	1–3 trips/year	Two negotiation units rather than four individuals; decisions made pairwise then merged
Multi-generation family	5–10	~1/year	Wide age and ability range; one or two adults do everything; constraints dominate preferences
Decisive small group	2–3	Frequent weekends	Short planning cycles, high trust, minimal tooling, fast commitment
Indecisive large group	6–10	Aspirational	Many contributors, high enthusiasm, plans frequently outlive their momentum
Once-a-year holiday group	3–6	1/year	High stakes per trip, budget-sensitive, planning starts months out

Roles within groups

- **The organiser.** Initiates, consolidates, chases, and often ultimately books. A striking share of survey respondents identified a single stable organiser in their group — frequently themselves — and the role appears sticky across trips rather than rotating.
- **The contributor.** Posts links and ideas enthusiastically, engages with others' suggestions, but does not consolidate or chase.

- **The "I'll go wherever" friend.** Genuinely flexible or conflict-avoidant; hard to distinguish which. Adds warmth to the thread and nothing to the decision.
- **The budget-conscious member.** Tracks cost implications of every option, sometimes silently; often the source of late-stage vetoes.
- **The constraint-holder.** Has fixed dates, dietary, mobility or childcare constraints that shape the feasible set more than anyone's preferences do.

Two cautions: roles are per-group, not per-person — the same individual can be the organiser in one group and the wherever-friend in another — and the archetypes describe planning behaviour, not demographics. Concept work should be stress-tested against at least the university friend group, the multi-generation family, and the decisive small group, which pull the design in different directions (breadth of contribution, asymmetry of ability, and speed).

6. Planning behaviours

This section walks the six journey stages, recording the behaviours the research has observed at each; the journey map artefact carries the sticky-level detail.

Discover

Discovery is abundant and effortless. Participants describe a constant background stream of trip-relevant material: social video, recommendations, articles, prior trips. Discovery is individual — one person's phone, one person's feed — and emotionally rewarding in itself; several participants collect places for hypothetical trips with no group and no date. Discovery does not stop when planning starts: new options keep arriving deep into Compare, re-opening what the group thought was settled.

Share

Sharing is immediate and low-ceremony: see something, forward it to the thread, add "this??" . The share sheet makes this a two-tap act, and groups do it constantly — our survey's sharing-frequency measures are among its highest-scoring behavioural items. Formats are heterogeneous: links that unfurl with a preview, links that unfurl badly, screenshots with no link, and occasionally a place name typed from memory. The format a place arrives in largely determines how much work it takes to do anything with it later.

Save

Saving is where individual effort begins and group participation thins out. The thread does not save things; a person does — the consolidation behaviour described in section 3 sits here. Personal maps-app lists also sit here, and in our data they are conspicuously private: people save places for the group trip into lists the group cannot see. Saving is also lossy and duplicative: the same place often exists as a chat message, a screenshot, a note entry and a personal pin, with no connection between the four.

Compare

Comparison is the most improvised stage. Comparing options needs them in one place, in one format, with the relevant attributes visible — and no tool in the current stack provides that. Groups improvise: side-by-side screenshots, spreadsheets, one member summarising five options in a long message that immediately scrolls away. Attributes must be re-looked-up repeatedly, and the answers live in whoever looked them up.

Participants describe this stage as effortful but not always as broken; several regard the improvisation as simply what planning is.

Decide

Deciding is the least tooled and least visible stage. Most groups have no artefact for it: no vote, no explicit moment of commitment, no record of what was agreed beyond the booking confirmations. Reported decision methods vary widely across the survey — group discussion, one person choosing, deference to whoever books, hybrids — and are heterogeneous enough that this pack refrains from characterising them further; read the survey's decision-method column and the interview transcript and form your own view.

Navigate

On the trip, planning artefacts are consulted less than expected. The organiser's note is referenced for the big fixed points; beyond that, groups improvise with fresh map searches. Saved places from planning are used for navigation only when the person navigating is the person who saved them. The trip generates its own new material, which feeds the next cycle of Discover.

7. Collaboration problems

The research so far surfaces recurring problems, reported here as participants expressed them. Frequency labels are relative to our small samples; the underlying rows and stickies are in the artefacts.

Problem (as reported)	Where reported	Relative frequency
"Things get lost in the chat" — links and ideas scroll away and are never found again	Survey free text, workshop, journey map	High
Keeping track of what has been shared across multiple apps and formats	Survey free text, interview	High
One person does all the work of consolidating and chasing	Survey, interview, journey map	High
The same places get re-shared and re-discussed multiple times	Journey map, interview	Medium
Screenshots arrive with no link, so places must be re-found by hand	Interview, journey map	Medium
The plan-of-record note goes stale and versions multiply	Journey map	Medium
Nobody responds when input is requested	Survey free text, interview	Medium
Dates drift because availability is never pinned down	Survey free text	Medium
Cost-splitting worries surface late and unsettle agreed plans	Survey free text, workshop	Low–medium
Group members use different phone platforms, limiting shared tooling	Survey free text	Low

When survey respondents were asked directly for their biggest planning frustration, the most common themes in the free text concerned sharing and keeping track — links getting lost, information scattered across apps, difficulty maintaining one view of what the group has found. This is consistent with the framing of the earlier studies summarised in section 17, and it is the framing the workshop's ideation largely followed.

A note of method rather than interpretation: the table mixes problems people **state** when asked with problems **observed** in their behaviour. The survey's behavioural columns, the interview's task observations and the telemetry in section 12 record what people do; the free text records what they say. Both are evidence. They are not the same evidence.

8. Ecosystem considerations (generic)

Any collaborative planning feature lives inside a mesh of other software, and the mesh shapes what is possible more than any single design decision. These considerations are generic to the category; none refers to a specific product or internal system.

Messaging apps are the incumbent venue, not a competitor. Every group we studied lives in a messaging thread, and no realistic outcome moves them out of it. The practical question is the interface between the thread and any planning surface: how a place travels from chat into a shared plan, and how the plan surfaces back into chat as something current and glanceable rather than a static title. A shared plan whose link unfurls as a rich, live preview inside the thread may matter more to adoption than most in-app features.

Cross-platform reality. Informal groups are mixed-platform by default. A shared plan that requires a particular device or app will, in most groups, exclude someone — disproportionately a peripheral contributor who then disengages entirely. Web fallbacks, read-only guest access, and graceful degradation for non-users are requirements, not enhancements.

Calendars, notes and documents adjoin the space. Dates live in calendar apps; consolidated plans live in notes and documents. The exploration does not aim to replace either, but should decide what it imports (dates, availability) and what it exports (a final plan as a shareable document). Every import/export boundary is also a failure point.

Booking is external and will remain so. Flights, trains, accommodation and restaurants book on third-party surfaces, so "booked" is a state the product can rarely observe directly; if it matters to the experience, it must be captured through user assertion or link handling, both unreliable in different ways.

Devices beyond the phone. Planning is phone-first, but comparison benefits from larger screens, and the Navigate stage implicates wearables and in-car contexts where glanceability and offline behaviour dominate. Noted for completeness; the first prototype is phone-only.

9. Accessibility considerations

Collaborative surfaces raise accessibility questions beyond single-user features: one member's inaccessible experience degrades the whole group's plan. These are design obligations for any concept.

Screen-reader coherence for shared state. A shared list is a live, multi-author surface: places appear and change state because of other people's actions, and unannounced remote changes make the surface untrustworthy for a screen-reader user. Concepts must define which changes are announced, how a member reviews "what changed since I last looked" linearly, and how attribution is exposed non-visually.

Status must never be colour-only. Any state we attach to a place — shortlisted, preferred, chosen, booked — must be legible through shape, label or position as well as colour, for colour-vision-deficient members and for the bright-sun and low-brightness contexts travel guarantees.

Dynamic type and layout resilience. Comparison layouts (cards, tables, side-by-side views) break first at large text sizes. Sketch concepts at accessibility text sizes from the start, not as a retrofit.

Motor accessibility. Map gestures (precise pans, long-presses, pin taps) are among the hardest for users with motor impairments; every map-mediated action needs a list-mediated equivalent.

Cognitive load and memory. Group planning spans weeks and members return after long gaps. Minimise reliance on remembering prior state — recaps, stable ordering and "where were we" affordances help

everyone, and are essential for members with memory or attention impairments.

Situational impairment while travelling. The Navigate stage is used one-handed, in noise, in sun, in a hurry, sometimes offline. Voice interaction, large targets and offline-tolerant behaviour are accessibility considerations here as much as conveniences.

Multi-generation groups make this concrete. The family archetype routinely spans a teenager and a grandparent on one plan. Accessibility is the difference between a plan the whole family shares and a plan two members operate on behalf of everyone — which quietly recreates the organiser burden the product hoped to relieve.

10. Product constraints

The exploration is open in its thinking but bounded in what it can ship. Constraints as currently understood:

- **Build on the existing saved-places model.** Shared trip plans must extend the current list/place architecture, not introduce a parallel place concept. A place in a shared plan is the same object a user saves privately.
- **No new social graph.** The product will not build friend lists, profiles or follower mechanics. Membership is established per-plan through invitations, and a plan's membership is its entire social universe.
- **No accounts requirement for viewing.** Per current direction, a shared plan should be viewable — though not editable — without an account, to serve mixed-platform groups.
- **Booking stays out of scope.** No booking engine, payment handling or cost-splitting ledger in v1. The product may record that something has been booked; it will not perform the booking.
- **Itineraries are out of scope for v1.** Day-by-day scheduling, routing and timetable generation are deferred. The v1 container is a plan (a set of places and their group state), not an itinerary.
- **Notification budget.** The maps product sends very few notifications today, and the bar for adding more is high. Assume a small number of high-value notifications per plan per week as a working ceiling.
- **Naming and language.** Alex flags that terms like "vote", "veto" and "deadline" carry social weight; terminology for group actions will go through content design review, and concepts should treat their labels as provisional.
- **Performance.** Shared lists must remain usable on mid-range hardware and poor connections; Compare must not require loading full place pages for every option.

None of these is immovable, but each has an owner and a cost to challenge. Concepts that respect them travel faster; those that break them must earn it.

11. Privacy considerations

For this exploration, privacy is a set of build and design constraints that any shared-planning concept must satisfy, deriving from the product's data-handling principles and the review every sharing feature undergoes before ship. They are stated here as requirements, not research findings.

- **Private by default.** Every list and saved place starts private to its creator. Nothing becomes visible to anyone else without an explicit act of sharing; there is no ambient visibility of any kind.
- **Sharing scope is per-plan and explicit.** Sharing a place into a group plan shares that place's entry in that plan — nothing else. It must not expose the sharer's other lists, saved-place history, home or work

locations, location history, or past trips. The design must make the boundary of what is shared legible at the moment of sharing.

- **Membership is enumerable and visible.** Any member of a shared plan can always see exactly who else can see it. Invite links must be revocable, and revocation must actually cut off access rather than merely hiding the link.
- **Contribution attribution is a design decision, not a default.** Whether "who added this place" and "who marked this preference" are visible to the group is a per-surface choice, made deliberately and disclosed in the interface. Attributed and unattributed models are both permissible; silent attribution is not.
- **Leaving is clean.** A member who leaves a plan can take their contributions with them or leave them behind — the design must choose and disclose which — and after leaving they retain access to nothing.
- **Data minimisation on the server.** Shared-plan infrastructure should store the minimum needed to synchronise state: place references, member state and plan metadata. It should not accumulate derived behavioural profiles of members, and plan data should have a defined retention and deletion story, including when the last member deletes the plan.
- **No location sharing.** A trip plan shares intentions, not positions. Live-location features are out of scope, and nothing in the plan should imply or reveal where any member currently is.
- **Age-appropriate defaults.** Multi-generation groups may include minors; invitation and visibility defaults must be safe for that case without special configuration.

These constraints interact with the concept space in non-obvious ways — unattributed preference-marking changes the social dynamics of choosing; explicit membership visibility constrains any "open invite" growth mechanic. Treat them as load-bearing from the first sketch, because retrofitting them reliably reshapes a sharing feature end to end.

12. Behavioural metrics (fictional)

The figures below come from a fictional instrumented sample of trip-style list behaviour over a recent 90-day window, indexed and rounded. "Trip-style lists" are personal or shared lists heuristically classified as trip-related (name patterns, place-geography spread, creation bursts). Definitions follow the table. These are observational measures: they record what happened, not why.

Metric	Value	Movement vs prior 90 days
Trip-style lists created (indexed)	100	+14%
Trip-style lists shared with at least one other person	58%	+6 pts
Share events per shared list (median)	9	+1
Place links sent outward to messaging apps, per trip-style list (median)	14	+2
Shared lists with two or more contributing members	24%	+2 pts
Multi-contributor lists reaching 10+ saved places	76%	+3 pts
Multi-contributor lists where any place is marked chosen or booked	31%	−1 pt
Median days from list creation to last edit	19	0
Lists reopened within 7 days of an inferred trip date	47%	+1 pt
Navigation sessions launched from a trip-style list	22% of lists	+2 pts

Definitions and caveats:

- **Shared**: the list owner generated a share link or added a member, and at least one other identity opened it. Sharing here measures access, not activity.
- **Contributing member**: an identity other than the owner that added, annotated or reordered at least one place. The gap between the 58% share rate and the 24% co-contribution rate means most shared lists are read-only in practice, whatever their permissions allow.
- **Share events**: outbound shares of the list or of individual places from it, counted per list. Alongside the outbound place-link figure, these indicate that sharing activity around trip-style lists is frequent and rising.
- **Chosen or booked**: any place in the list was flagged with the favourite/priority marker, or a booking-type outbound link (reservation, tickets, accommodation) was tapped from a place in the list. This is a permissive, best-effort proxy; it will overcount casual flagging and undercount bookings completed from outside the list.
- The step from 76% (multi-contributor lists reaching 10+ saved places) to 31% (any place marked chosen or booked) is the largest single drop in the funnel. Current instrumentation records where lists stop progressing but contains nothing that explains it; no event we log distinguishes a list that quietly served its purpose from one that did not.
- The trip-date inference used in the reopen metric is weak (calendar-adjacent signals only) and should not be leaned on.
- Sample is not longitudinal per group; a group planning two trips appears twice.

Taken together: list creation is healthy and growing, sharing is frequent, genuine co-editing is markedly rarer than sharing, and lists accumulate places far more reliably than they arrive at marked choices.

13. Existing assumptions

The exploration inherits assumptions from earlier work and from the team's own instincts. The register is maintained by Nina and reviewed fortnightly; statuses reflect evidence held at this pack's issue date. Presence here is not endorsement — the register exists so inherited beliefs are visible enough to be tested.

ID	Assumption	Origin	Status	Notes
A1	Users lack an easy way to share places with each other	Saved Places Foundations study (2024)	Untested — carried forward	Headline recommendation of the 2024 study; adopted into this exploration's framing at kick-off and not re-examined since
A2	Group trip planning is done primarily on phones	Group Planning Survey (2026)	Supported	36 of 40 respondents phone-first; larger screens appear mainly at comparison
A3	Each group has one stable organiser who does most of the work	Survey + Interview 01	Supported (directional)	Strong signal in both sources; sample remains small
A4	People want their personal saved lists kept separate from anything the group can see	Interview 01	Directional — single participant	One participant's clear preference; needs wider validation before it constrains architecture
A5	Larger groups take longer to reach agreement than smaller groups	Team belief	Untested	Widely assumed in the workshop; no behavioural data either way

ID	Assumption	Origin	Status	Notes
A6	Groups will not adopt a planning surface that requires leaving their group chat	Workshop discussion	Contested	The workshop split; ecosystem reasoning (section 8) supports taking it seriously, but no direct evidence exists
A7	Booking is completed outside the maps product and will remain so	Telemetry + survey	Supported	Consistent across every source; also a scoping decision (section 10)
A8	A single shared list per trip is the right container for group planning	Trip Spaces concept	Untested	Baked into the current prototype without examination; alternative containers unexplored

Assumptions A1, A5 and A8 are the most deeply embedded in current concept work relative to the evidence behind them. When reviewing the Trip Spaces prototype and the workshop output, ask which register entry each feature silently depends on.

14. Hypotheses

The team maintains falsifiable hypotheses to keep concept work testable; each states an expected effect and what would count as disconfirmation. None has been tested yet; the planned vehicle is a multi-group diary-plus-prototype study in weeks 7–9.

- **H1 — Consolidation.** If a group has a single shared space that captures places as they are shared, duplicate re-shares of the same place in chat will fall substantially (working threshold: 30% over a planning cycle). *Disconfirmed if* groups re-share at similar rates, or route around the space and keep consolidating in notes.
- **H2 — Structure accelerates decisions.** Groups given lightweight decision structure — a shortlist distinct from the pool, per-member preference marks, and an explicit "locked" state — will reach a committed choice in fewer days than matched groups using an unstructured shared list. *Disconfirmed if* time-to-commitment is unchanged, or if structured groups abandon the structure before reaching it.
- **H3 — Visible state broadens participation.** If the plan makes each member's activity visible (who has added, who has responded), more members will perform at least one save or preference action. *Disconfirmed if* contribution remains concentrated in the organiser at current levels.
- **H4 — Comparable cards reduce re-work.** Presenting options as uniform cards with group-relevant attributes (hours, price band, distance) will reduce repeated look-ups of the same place during comparison. *Disconfirmed if* participants still open full place pages or external sources at baseline rates.
- **H5 — Explicit commitment carries forward.** Plans containing at least one explicitly "locked" place will show higher downstream engagement — reopens near the trip date and navigation sessions launched from the plan — than plans without. *Disconfirmed if* locked and unlocked plans show similar downstream behaviour, suggesting locking records commitment rather than creating it.

H2 and H5 are the most load-bearing for the current prototype direction; H2 is also the most expensive to test, since it requires observing whole groups. Study design for it is an open task Maya would welcome help with.

15. Success measures

If a concept shipped, these measures would tell us whether it worked, organised against the journey and paired with guardrails — a collaborative feature can succeed on engagement while failing the group it serves.

Measure	Journey stage	Type	Working target
Places entering a shared plan per group per week	Share/Save	Leading	Above current per-person save rates for trip-style lists
Share of group members who contribute at least once	Save/Compare	Leading	Majority of members in a majority of plans
Shared plans in which any place reaches a chosen/booked state	Decide	Lagging	Materially above the current 31% baseline (section 12)
Median days from plan creation to first chosen place	Decide	Lagging	Down, relative to observed planning cycles
Plan reopens in the week before the trip	Navigate	Lagging	Above the current 47% list-reopen baseline
Navigation sessions launched from the plan during the trip	Navigate	Lagging	Above the current 22% baseline
Organiser-reported effort (pulse survey)	All	Qualitative	Down vs pre-launch cohort
Group-reported satisfaction with the outcome (pulse survey)	All	Qualitative	Up vs pre-launch cohort

Guardrails:

- **Notification burden:** notifications per member per week stay within the section 10 ceiling; mute/leave rates stay low.
- **Privacy comprehension:** in usability testing, members can correctly state who can see a plan and what sharing a place exposes; failure here blocks ship regardless of other results.
- **No participation theatre:** contribution breadth must reflect real input, not gamified tapping; H3-style metrics are read alongside qualitative evidence.
- **Organiser displacement, not organiser addition:** if the feature adds coordination work on top of the chat rather than absorbing any, organiser effort scores will show it, and that counts as failure even with healthy usage.

Several lagging measures depend on the permissive chosen/booked proxy described in section 12; before any launch, instrumentation for commitment states needs tightening, or we will be unable to distinguish success from noise.

16. Open questions

Questions the current evidence cannot answer; each names what would resolve it.

1. **What is the right container?** One plan per trip is the prototype's guess (assumption A8); per-day, per-city or per-category structures are unexplored. *Resolved by:* concept variants in the weeks 7–9 study.
2. **Should the product assume an organiser exists?** Designs can lean into the organiser role or deliberately flatten it. Which does a real group actually sustain? *Resolved by:* multi-group observation;

single-participant data cannot answer this.

3. **How should disagreement surface?** Preference marks, comments, vetoes and polls all externalise disagreement differently, and the social cost of each is unknown. *Resolved by:* group sessions; attitudes to voting-style mechanics vary sharply in our sources and need direct observation.

4. **Would deadlines or reminders help groups decide?** The workshop proposed date-anchored nudges ("flights at this price until Friday", "you said you'd decide by Sunday"), but the survey never asked about them, no behavioural data touches them, and no participant has been shown one. Whether time pressure unblocks groups or reads as nagging from a maps app is genuinely unknown. *Resolved by:* prototype variants with and without time-anchored prompts.

5. **Where do groups enter?** From a chat link, the maps app, or a shared place? The entry point shapes adoption more than any feature, and we have no evidence which is natural. *Resolved by:* diary study capture of first-touch moments.

6. **What happens after the trip?** Does a plan become a memento, a template for the next trip, or clutter? *Resolved by:* retrospective interviews with groups who travelled recently.

7. **How does the experience degrade for non-members?** Mixed-platform groups guarantee some members see the plan only as links and previews. Is a read-only web view enough to keep them in the group's process? *Resolved by:* testing with real mixed groups.

8. **Do multi-destination trips break the model?** Touring trips multiply anchors and distances; card comparison against a single anchor stops making sense. *Resolved by:* concept stress-testing with touring-trip planners.

9. **Should chosen/booked be asserted or inferred?** User assertion is honest but effortful; inference from booking-link taps is automatic but wrong in both directions. *Resolved by:* instrumentation prototyping and observing whether users assert.

17. Previous research summaries (fictional)

Three earlier studies inform this exploration's framing; these summaries cover scope and headline findings.

Saved Places Foundations (2024)

A three-week diary study with 18 participants examining how people save, organise and retrieve places in the maps product and elsewhere. Core findings: people save far more places than they ever visit, and derive real pleasure from the collecting itself — saved lists function as aspirational memory as much as practical planning; retrieval is weak, and duplicated saves across apps were near-universal. The study's most-cited finding concerned friction at the boundaries: getting a place from where it was discovered (a video, a message, a webpage) into a durable saved form took too many steps, and getting a saved place in front of another person took more. The report's headline recommendation — **make sharing places effortless** — has shaped saved-places roadmap thinking since, and is the origin of assumption A1 in this exploration's register.

Link Sharing Habits (2025)

A mixed-methods study (survey n=220 plus 10 interviews) of how place-related links travel between apps, prompted by the growth of place discovery in social video. Findings centred on capture and transport: many discoveries arrive as screenshots or app-specific links that strip metadata in transit, and participants

described "screenshot graveyards" in their photo libraries. The report mapped a "leaky pipeline" from discovery to saved place and recommended investment in capture and richer outbound sharing formats. Its framing — places are easy to find and hard to move — is widely referenced and was the second pillar of this exploration's original pitch.

Weekend Planners (2023)

Twelve interviews with frequent short-trip takers, focused on inspiration and discovery habits. The study predates the group focus — participants were interviewed as individuals — but documented behaviours that recur in our data: place collections for cities not yet visited, friend recommendations trusted above editorial sources, and sharing finds as social currency ("sending someone a place is like sending them a little gift"). Its recommendations concerned discovery surfaces and richer collecting tools; it remains the team's best source on the emotional texture of collecting.

Read together, the three studies build a consistent picture of the collecting and sharing life of places, and explain why this exploration began where it did. None of the three observed a group making a decision together; group decision-making was outside each study's scope.

18. Competitive and service patterns (generic)

A category-level scan of how other products approach group planning, described generically; no specific product or non-public information is referenced.

Board-style travel planners. A cluster of travel apps organise trips as visual boards: places as cards, grouped into sections, with members adding and commenting. They excel at collecting and the aesthetic pleasure of a growing plan; public reviews suggest groups enjoy filling boards. Commitment mechanics (choosing, booking states) are typically thin or absent.

Itinerary compilers. Products that assemble day-by-day itineraries, often importing confirmation emails automatically. Strong after decisions are made — everything booked, in order, with times — and weak before it. Their popularity suggests real demand for a "single source of truth" artefact, positioned after the stage this exploration studies most.

Document and spreadsheet templates. A striking amount of group planning happens in general-purpose docs and spreadsheets using community-circulated templates (comparison grids, budget tabs, voting columns). The pattern's persistence signals that flexible tooling has genuine pull — and that people will do surprising manual work inside a tool the whole group can access.

Group mechanics from adjacent domains. Patterns worth studying from outside travel: polls in messaging apps (low-friction, widely used for scheduling); event-RSVP products (explicit yes/no/maybe states and visible non-response); wishlist-and-claim mechanics in commerce; watch-list voting in streaming households. Each solves a narrow slice of group agreement with a distinct social contract around visibility and attribution.

Split-cost tools. Widely adopted for the money side of group travel; notable for succeeding with a strict, narrow model and for being adopted mid-trip, not during planning. Their lesson is scope discipline as much as features.

Pattern-level observations. Across the category: collecting surfaces are abundant and polished; deciding surfaces are rare and mostly borrowed from scheduling (polls, RSVPs); and no widely adopted product in the scan treats the transition from "many options" to "a committed choice" as its central design problem.

Whether that is a gap or a graveyard cannot be determined from public information; read the scan with that ambiguity intact.

19. Technical considerations

Sam's feasibility notes, summarised for design relevance. None blocks concept work; several shape what is cheap versus expensive to promise.

Shared state and conflict. A collaborative list is a synchronised structure edited concurrently on devices that are frequently offline. The infrastructure pattern is well understood — operation-based merging with eventual consistency — but its design consequences deserve early attention: two members can add the same place seconds apart (de-duplicate by place identity), ordering edits interleave unpredictably (avoid making ordering carry meaning), and "who did what when" must be recorded carefully if attribution is surfaced. Free-text notes are the hardest merge case; structured state merges cleanly.

Place identity over time. A plan lives for weeks to months; places change under it — venues close, hours change, listings disappear. Place references must be stable identifiers resolved to live data at view time, with a defined degraded state for places that no longer resolve, treated as normal for any plan older than a few weeks.

Invitation and access. Link-based invites are the adoption-friendly choice and the security-sensitive one: links get forwarded. The requirements in section 11 translate to tokenised links with server-side validity, per-member access records, and a real view/edit distinction — all standard, but they make "anyone with the link can edit" a configuration we would need to argue for explicitly.

Previews and unfurling. If a plan's chat-link preview is to be live and glanceable (section 8), preview rendering must be served without authentication while leaking nothing sensitive — likely a minimal summary (name, place count, cover image) rather than plan contents. Freshness across messaging apps' caches is only partially under our control.

Notifications and activity. Server-side activity aggregation is needed to honour the notification budget: batching ("3 new places this week") rather than per-event pings. The activity feed and notification stream should share one event model so the in-app recap absorbs what notifications do not send.

Attribute retrieval for comparison. Card-based comparison (H4) needs a subset of place attributes fetched in bulk cheaply — thinner than full place pages. Feasible, but fix the attribute set early; each addition has a data-quality tail (price-band coverage varies sharply by region and category).

Scale posture. Group sizes are small and edit rates low; this is not a throughput problem. Engineering cost centres on correctness (merge, access, deletion) and on plans' long idle lifetimes, which make retention and archival policy (section 11) a first-class requirement.

20. Risks and dependencies

Risk	Likelihood	Impact	Mitigation posture
Concept solves a stage groups were not struggling with, missing the real breakdown	Medium	High	The weeks 7–9 study tests stage-level value, not polish; synthesis must revisit which stage the evidence indicts
Social friction: making group state visible creates awkwardness that suppresses use	Medium	High	Test attributed vs unattributed variants; treat guardrail qualitative signals as blocking

Risk	Likelihood	Impact	Mitigation posture
Notification fatigue causes plan muting and abandonment	Medium	Medium	Hard budget (section 10); batching (section 19); mute metrics as guardrail
Mixed-platform groups exclude members, collapsing group adoption	High	Medium	Read-only web access in v1; test with genuinely mixed groups
Scope creep towards a full itinerary/booking product	High	Medium	Section 10 constraints have named owners; concept reviews check scope
Place-data quality undermines the comparison surface in some regions	Medium	Medium	Fix the card attribute set early; design degraded states as normal (section 19)
Privacy review requires late changes to sharing mechanics	Low	High	Section 11 constraints applied from first sketch; early consultation rather than end-stage review
Research base too thin: one interview and n=40 carry more framing weight than they should	High	Medium	Interviews 02–04 and the group study are the priority; treat single-source findings as provisional
Chosen/booked instrumentation proxy too weak to measure success	Medium	High	Instrumentation tightening scoped as a pre-launch dependency (section 15)
Team framing inherited from prior studies narrows the concept space prematurely	Medium	High	Assumptions register reviewed fortnightly; new-joiner reviews explicitly invited to challenge it

Dependencies. The weeks 7–9 multi-group study is the load-bearing dependency for nearly every open question and hypothesis; recruiting real, intact friend groups is its critical path and historically slow. Sync and access infrastructure has no committed engineering line yet — Sam's notes are feasibility, not a plan of record. Content design review of group-action terminology gates external testing. And the twelve-week window makes synthesis discipline itself a dependency: the evidence we already hold has to be mined thoroughly before the instinct to gather more takes over.

Your first contribution is exactly that mining: read the four artefacts end to end against this pack, and arrive at the week-six review with your own account of where — across Discover, Share, Save, Compare, Decide and Navigate — the strongest evidence says this product should intervene first.